

# Denoising Diffusion Probabilistic Models (DDPM) - A Complete Guide with PyTorch Implementation

GitHub: <https://github.com/musawir124/Denoising-Diffusion-Probabilistic-Models.git>

Diffusion models are a class of **generative models** that have gained **massive popularity** due to their ability to generate high-quality, diverse, and realistic images. These models **reverse a noise process**, gradually converting random noise into meaningful images.

This tutorial will guide you through the **theory and implementation** of DDPMs, explaining the mathematical foundations and coding the model in **PyTorch** step by step.

## 1. Introduction to Generative Models

Generative models are **AI models** that can create **new data** similar to a given dataset. Some popular generative models include:

- **GANs (Generative Adversarial Networks)**
- **VAEs (Variational Autoencoders)**
- **Flow-based Models**
- **Diffusion Models (DDPMs, Stable Diffusion, etc.)**

### 💡 Why use Diffusion Models?

- **Stable Training** compared to GANs.
- **High-Quality Image Generation.**
- **Better Diversity** (avoids mode collapse seen in GANs).

## 2. What is a Diffusion Model?

A **Diffusion Model** consists of two processes:

### ◆ Forward Process (Adding Noise)

A clean image is **gradually corrupted** by adding Gaussian noise at each time step until it becomes **pure noise**.

### ◆ Reverse Process (Denoising)

A neural network is trained to **undo the noise** step-by-step, reconstructing the original image.

Forward Process (Noise Addition):

The **goal** of the forward process is to **progressively destroy** an image by adding noise step by step.

At each step  $t$ , we add Gaussian noise:

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I)$$

where:

- $x_t$  is the noisy image at time step  $t$ .
- $\alpha_t$  controls the noise level.
- $I$  is the identity matrix, ensuring noise follows a Gaussian distribution.

By applying the **reparameterization trick**, we can directly compute  $x_t$  from  $x_0$  (original image):

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$$

where:

- $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$  (cumulative product of noise schedule).
- $\epsilon \sim \mathcal{N}(0, I)$  (random Gaussian noise).

Reverse Process (Denoising):

The reverse process **reconstructs** the original image from pure noise.

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma^2 I)$$

- A **neural network**  $\epsilon_\theta(x_t, t)$  is trained to predict noise at each step.
- The predicted noise is then used to gradually **denoise** the image.

Mathematical Formulation of DDPM:

**Loss Function:**

Instead of learning  $p(x_{t-1}|x_t)p(x_{t-1}|x_t)$  directly, we train a neural network to predict the noise  $\epsilon$  that was added at each step.

$$L_{\text{simple}} = \mathbb{E}_{x_0, \epsilon, t} [||\epsilon - \epsilon_{\theta}(x_t, t)||^2]$$

The model learns to predict the noise **accurately** so that we can subtract it and reconstruct the original image.

Let's code a basic DDPM from scratch in PyTorch!

```
In [8]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

```
In [9]: class SimpleDiffusionModel(nn.Module):
def __init__(self):
    super(SimpleDiffusionModel, self).__init__()
    self.encoder = nn.Sequential(
        nn.Conv2d(1, 32, 3, padding=1),
        nn.ReLU(),
        nn.Conv2d(32, 64, 3, padding=1),
        nn.ReLU()
    )
    self.decoder = nn.Sequential(
        nn.ConvTranspose2d(64, 32, 3, padding=1),
        nn.ReLU(),
        nn.ConvTranspose2d(32, 1, 3, padding=1),
        nn.Sigmoid()
    )

def forward(self, x):
    x = self.encoder(x)
    x = self.decoder(x)
    return x
```

```
In [10]: transform = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize((0.5,), (0.5,))
        ])

train_dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
train_loader = DataLoader(train_dataset, batch size=64, shuffle=True)
```

```
In [11]: def add_noise(images, noise_level=0.5):
          noise = torch.randn_like(images) * noise_level
          return images + noise
```

```
In [12]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = SimpleDiffusionModel().to(device)
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.MSELoss()
```

```

num_epochs = 5

for epoch in range(num_epochs):
    for images, _ in train_loader:
        images = images.to(device)
        noisy_images = add_noise(images)

        optimizer.zero_grad()
        output = model(noisy_images)

        loss = criterion(output, images)
        loss.backward()
        optimizer.step()

    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}")

# Save trained model
torch.save(model.state_dict(), "diffusion_model.pth")

```

Epoch [1/5], Loss: 0.8375

Epoch [2/5], Loss: 0.8345

Epoch [3/5], Loss: 0.8551

Epoch [4/5], Loss: 0.8354

Epoch [5/5], Loss: 0.8418

## Generate New Images

```


In [13]: def generate_images(model, noise_level=0.5):
        model.eval()
        noise = torch.randn((1, 1, 28, 28)).to(device)
        with torch.no_grad():
            denoised_image = model(noise)
        return denoised_image

# Load Model
model.load_state_dict(torch.load("diffusion_model.pth"))
model.to(device)

# Generate Image
generated_img = generate_images(model).cpu().squeeze().numpy()

plt.imshow(generated_img, cmap="gray")
plt.show()

```



## Final Summary

Defined a simple U-Net model for denoising.

Added noise to images and trained the model.

Saved the trained model for future use.

Generated new images by reversing the noise process.

